

Lab 12 Unauthorized access investigation

AI for IT Security Professionals | TGS-2023039344 | v6.0

Purpose and output

Create a chronological evidence timeline and incident brief. This lab supports LO4. It uses a simplified deterministic teaching model, not a trained AI system or full Azure emulator. The AI review is a separate exercise using the included prompt.

Requirements

Python 3.10 or later, a text editor and this entire folder. No packages, network access or API key are required for the baseline. On Windows use `py -3` if `python3` is unavailable. An approved AI tool is optional; the trainer can provide a review response for critique. Azure exercises require the trainer-assigned subscription, permission and feature entitlement. Record an exact blocker where unavailable; never describe an unperformed test as passed.

Folder contents

- `mock-data.json`: synthetic input with no real personal data or credentials.
- `analyse.py`: runnable analysis for this lab only.
- `expected-results.json`: expected baseline and source checksum.
- `ai-review-prompt.txt`: evidence-constrained AI review prompt.
- `evidence-template.md`: your deliverable template.
- `INSTRUCTIONS.md` and `INSTRUCTIONS.pdf`: the same procedure in two formats.

1 Prepare a working copy

Copy this whole folder to a location you can edit. Open a terminal in the copied folder. Confirm Python and inspect the data before running code.

```
python3 --version
python3 -m json.tool mock-data.json
```

Identify the record IDs and explain the meaning of each field. All values are invented classroom fixtures. Open `analyse.py` in the editor and locate the decision rule. It reads a local JSON file and writes a local report.

2 Run the baseline

```
python3 analyse.py --input mock-data.json --output results.json
```

Open `results.json`. The expected decisions in output order are: `signin`, `role-grant`, `read-data`. The `source_sha256` binds the report to the exact input bytes. Compare against the supplied baseline:

```
python3 -c "import json; a=json.load(open('results.json')); b=json.load(open('expected-results.json')); assert a==b; print('BASELINE PASS')"
```

A pass proves this fixture was processed as specified; it does not prove an Azure control was deployed. Copy the input checksum and decision rows into your evidence template.

3 Trace the mechanism

Evidence led investigation

A hypothesis is tested against events rather than treated as a fact because AI repeats it.

Contract: hypothesis, supporting_ids, contradicting_ids, status

Accepted case: An access event and role change support a privilege-use hypothesis.

Counterexample: Only a high anomaly score supports a claimed compromise.

Diagnosis: Collect corroboration and record uncertainty before containment decisions.

Evidence: Maintain a hypothesis table with supported, contradicted and unresolved states.

Incident timeline

Order events by event time while retaining sources and integrity hashes.

Contract: timestamp_utc, event_id, actor, source_hash

Accepted case: E03 role grant precedes E07 resource read.

Counterexample: A summary reverses the order and claims the read caused the grant.

Diagnosis: Rebuild chronology from raw evidence and preserve original files unchanged.

Evidence: The timeline must cite exact event IDs and UTC timestamps.

Containment approval

The analyst chooses a proportionate action and an authorised owner approves its execution.

Contract: incident_id, action, approver, rollback, status

Accepted case: A test account suspension is approved with a recovery owner.

Counterexample: An AI summary directly disables a production administrator.

Diagnosis: Hold execution until approval and scope checks pass; preserve forensic evidence first.

Evidence: Record decision rationale, authority, action result and recovery route.

4 Change one input and test the consequence

Save a copy of mock-data.json as changed-data.json using the editor. Change E07 minute to 2. The read now precedes the role grant. Revise the hypothesis instead of retaining the original causal claim.

```
python3 analyse.py --input changed-data.json --output changed-results.json
```

Compare decisions and source hashes with the baseline. Identify the exact expression in analyse.py responsible for the changed behaviour. Do not overwrite expected-results.json to make a failing baseline pass. Restore the original fixture if you edited it accidentally.

5 Review AI claims against evidence

Open ai-review-prompt.txt. In an organisation-approved AI tool, submit the prompt with the synthetic input and result only. Record the tool/model name, date and prompt version if available. If no approved AI tool is available, manually draft a tentative analyst summary and label it as a human exercise.

For each returned claim, record its cited ID, the actual field value, whether the claim follows, and any correction.

Reject conclusions unsupported by the records even when the model is confident. Include one alternative explanation. The AI response is a proposal; it has no authority to approve or execute a security action.

6 Verify the corresponding operational control

Microsoft Sentinel > Incidents > trainer-provided incident. Inspect entities, alerts and source events. Record UTC timestamps and correlation IDs. Compare at least two hypotheses and identify missing evidence. Containment requires a named authorised decision; do not disable real accounts as part of this exercise.

For every screenshot or export, record resource scope, UTC time and expected versus observed result. Redact personal data and secrets. If blocked, record the exact error, missing permission/licence, what could be inspected, and the unverified outcome. Ask the trainer to provide an authorised sandbox demonstration where the assessed ability cannot otherwise be evidenced.

7 Submit evidence and clean up

Complete evidence-template.md with baseline, changed result, AI claim review and Azure evidence or clearly labelled limitations. Keep results.json and changed-results.json with your submission. Check that your conclusion distinguishes an observed fact from a hypothesis. Remove only resources or assignments you created with trainer approval; do not delete shared training infrastructure.

Acceptance checks

- The unchanged fixture produces the exact expected report.

- The changed fixture produces the explained result and a different input hash.
- At least one AI or analyst claim is checked against a specific record and field.
- The report states simulation, Azure verified or blocked for each relevant claim.
- The output is a chronological evidence timeline and incident brief with an owner, evidence and remaining uncertainty.

Troubleshooting

- Python not found: install the trainer-approved Python distribution or use `py -3` on Windows.
- File not found: open the terminal in this lab folder; check the input filename.
- JSON parse error: validate with `python3 -m json.tool changed-data.json`; use lowercase `true` and `false` in JSON.
- Baseline mismatch: restore the supplied fixture and inspect the first differing decision; never edit the expected file.
- Azure denial or missing feature: capture the exact error and scope. A blocker is a limitation, not proof of competence or deployment.
- AI invents evidence: reject the claim, cite the missing ID, and request an evidence-limited revision.

References

The Learner Guide contains the complete source register and the lecture explanations for this lab. Original examples are adapted from the supplied Azure and AI-security references. Product behaviour should be checked against current Microsoft Learn documentation before a real deployment.